

▼ 1 Intro: Basic Curve Fitting with Gradient Descent

▼ Setting

Inputs

```
1 import random
2 import math
3
4 n = 20
5 alpha = 0.01
6 epsilon = 0.1 ** 5
```

▼ Make initial dataset

```
1 random.seed(42)
2 x = []
3 y = []
4 for i in range(n):
5     x.append(random.random())
6     y.append(random.random())
7 print("x: ", x)
8 print("y: ", y)
```

x: [0.6394267984578837, 0.27502931836911926, 0.7364712141640124, 0.8921795677048454, 0.4219218196852704, 0.21863797480360336, 0.0265359696838636, 0.025010755222666936, 0.22321073814882275, 0.6766994874229113, 0.08693883262941615, 0.029797219438070344, 0.5053552881033624, 0.198837650686, 0.10914296454641485, 0.004173947614718544, 0.8599022101180511, 0.15589485167486819, 0.8456462542360882]

y: [0.025010755222666936, 0.22321073814882275, 0.6766994874229113, 0.08693883262941615, 0.029797219438070344, 0.5053552881033624, 0.198837650686, 0.10914296454641485, 0.004173947614718544, 0.8599022101180511, 0.8456462542360882, 0.025010755222666936, 0.22321073814882275, 0.6766994874229113, 0.08693883262941615, 0.029797219438070344, 0.5053552881033624, 0.198837650686, 0.10914296454641485, 0.004173947614718544]

▼ Guess m and b

```
1 m = 0.5
2 b = 0
3 print("m: ", m, " b: ", b)
```

m: 0.5 b: 0

▼ Gradient Descent Process

```
1 detected = False
2 while not detected:
3     delta_m = 0
4     delta_b = 0
5     for i in range(n):
6         delta_m += (2*b*x[i] + 2*m*(x[i] ** 2) - 2*x[i]*y[i])
7         delta_b += (2*b + 2*m*x[i] - 2*y[i])
8     delta_m = delta_m / n
9     delta_b = delta_b / n
10    m -= alpha * delta_m
11    b -= alpha * delta_b
12    if math.sqrt(delta_m**2 + delta_b**2) < epsilon:
13        detected = True
14
15 print("m: ", m, " b: ", b)
```

m: 0.23606845773749743 b: 0.27746866427140193

▼ Scikit-learn Linear Regression

```

1 from sklearn.linear_model import LinearRegression
2 import pandas as pd
3 import matplotlib.pyplot as plt
4
5 mod = LinearRegression()
6 df_x = pd.DataFrame(x)
7 df_y = pd.DataFrame(y)
8 mod_lin = mod.fit(df_x, df_y)
9 print('m:', mod.coef_[0])
10 print('b:', mod.intercept_)

```

```

m: [0.2359993]
b: [0.27751048]

```

We can confirm that the answers are almost the same in both processes.

▼ 2 Building and Training a Neural Network for Rasterized Digit Classification

▼ Setup

▼ Load Data

```

1 # Reference: https://github.com/sorki/python-mnist/blob/master/mnist/loader.py
2 import struct
3 import numpy as np
4
5 train_image_path = './data/train-images-idx3-ubyte'
6 train_label_path = './data/train-labels-idx1-ubyte'
7 test_image_path = './data/t10k-images-idx3-ubyte'
8 test_label_path = './data/t10k-labels-idx1-ubyte'
9 with open(train_label_path, 'rb') as file:
10     magic, size = struct.unpack(">II", file.read(8)) # 使わないけど読み込んだかないとlabelsに余計なものが読み込まれちゃう。
11     labels = file.read()
12 with open(train_image_path, 'rb') as file:
13     magic, size, rows, cols = struct.unpack(">IIII", file.read(16))
14     image_data = file.read()
15 print("size:", size, ' rows:', rows, ' cols:', cols)
16 images = []
17 for i in range(size):
18     images.append([0] * rows * cols)
19     images[i][:] = image_data[i * rows * cols:(i + 1) * rows * cols]

```

size: 60000 rows: 28 cols: 28

```

1 with open(test_label_path, 'rb') as file:
2     magic, size_test = struct.unpack(">II", file.read(8))
3     labels_test = file.read()
4 with open(test_image_path, 'rb') as file:
5     magic, size_test, rows, cols = struct.unpack(">IIII", file.read(16))
6     image_data_test = file.read()
7 print("size:", size_test, ' rows:', rows, ' cols:', cols)
8 images_test = []
9 for i in range(size_test):
10     images_test.append([0] * rows * cols)
11     images_test[i][:] = image_data_test[i * rows * cols:(i + 1) * rows * cols]

```

size: 10000 rows: 28 cols: 28

```

1 train = []
2 for i in range(size):
3     label_onehot = [0] * 10
4     label_onehot[labels[i]] = 1
5     train.append([np.reshape(images[i], (-1,1)), np.reshape(label_onehot, (-1,1))])
6 test = []
7 for i in range(size_test):
8     label_onehot = [0] * 10
9     label_onehot[labels_test[i]] = 1
10     test.append([np.reshape(images_test[i], (-1,1)), np.reshape(label_onehot, (-1,1))])
11

```

▼ Show Image

```

1 # As numbers
2 img_num = 1
3 for idx, val in enumerate(images[img_num]):
4     if idx%28==0:
5         print('\n')
6         print(str(val).center(3), end=' ')
7 print('\n\nLabel: ', labels[img_num])

```

```

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 51 159 253 159 50 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 48 238 252 252 252 237 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 54 227 253 252 239 233 252 57 6 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 10 60 224 252 253 252 202 84 252 253 122 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 163 252 252 252 253 252 252 96 189 253 167 0 0 0 0 0
0 0 0 0 0 0 0 0 0 51 238 253 253 190 114 253 228 47 79 255 168 0 0 0 0 0
0 0 0 0 0 0 0 48 238 252 252 179 12 75 121 21 0 0 253 243 50 0 0 0 0 0
0 0 0 0 0 0 38 165 253 233 208 84 0 0 0 0 0 0 253 252 165 0 0 0 0 0
0 0 0 0 0 7 178 252 240 71 19 28 0 0 0 0 0 0 253 252 195 0 0 0 0 0
0 0 0 0 57 252 252 63 0 0 0 0 0 0 0 0 0 253 252 195 0 0 0 0 0
0 0 0 198 253 190 0 0 0 0 0 0 0 0 0 0 255 253 196 0 0 0 0 0
0 76 246 252 112 0 0 0 0 0 0 0 0 0 0 253 252 148 0 0 0 0 0
0 85 252 230 25 0 0 0 0 0 0 0 0 7 135 253 186 12 0 0 0 0 0
0 85 252 223 0 0 0 0 0 0 0 0 7 131 252 225 71 0 0 0 0 0
0 85 252 145 0 0 0 0 0 0 0 48 165 252 173 0 0 0 0 0 0 0
0 86 253 225 0 0 0 0 0 0 114 238 253 162 0 0 0 0 0 0 0
0 85 252 249 146 48 29 85 178 225 253 223 167 56 0 0 0 0 0 0
0 85 252 252 252 229 215 252 252 252 196 130 0 0 0 0 0 0 0
0 28 199 252 252 253 252 252 233 145 0 0 0 0 0 0 0 0 0 0
0 25 128 252 253 252 141 37 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

```

```


1 import matplotlib.pyplot as plt
2 import numpy as np
3
4 def mnist_digit_show(flatimage, outname=None):
5     image = np.reshape(flatimage, (28, 28))
6     plt.matshow(image, cmap=plt.cm.binary)
7     plt.xticks([])
8     plt.yticks([])
9     if outname:
10         plt.savefig(outname)
11     else:
12         plt.show()

```

```

1 mnist_digit_show(images[0])

```



▼ Data for debugging

```
1 import pickle
2 import gzip
3
4 location = './debugdata/tinyMNIST.pkl.gz'
5 f = gzip.open(location, 'rb')
6 u = pickle._Unpickler(f)
7 u.encoding = 'latin1'
8 train, test = u.load()
```

▼ Neural Network

▼ Inputs

- nl: number of dense (fully connected) linear layers in your NN (excluding the first and last layers)
- nh: number of units in each of the hidden layers
- ne: number of training epochs
- nb: number of training samples per batch
- alpha: learning rate
- lam: Regularization parameter.

```
1 # nl = 2
2 # nh = 30
3 # ne = 10000
4 # nb = 6000
5 # alpha = 0.1
6 # lam = 0.0001
7 # output_type = 0
8
9 nl = 2
10 nh = 30
11 ne = 20000
12 nb = 10
13 alpha = 0.1
14 lam = 0.0001
15 output_type = 0
```

▼ Functions

```
1 import random
2 np.random.seed(42)
3 random.seed(42)
4
5 layers = [train[0][0].shape[0]] + [nh]*nl + [train[0][1].shape[0]]
6 L = len(layers)
7 # biases = [np.random.randn(n, 1) for n in layers[1:]] # input layer以外で必要. sizeが((30,1),(30,1),(10,1))のlist of lists.
8 biases = [np.zeros((n, 1)) for n in layers[1:]] # input layer以外で必要. sizeが((30,1),(30,1),(10,1))のlist of lists.
9 weights = [np.random.randn(n, m) for (m, n) in zip(layers[:-1], layers[1:])] # ((784,30),(30,30),(30,10)). 長さmの縦長行列を長さnに変換する。
10 deltas = [[np.zeros((n, 1)) for n in layers]] * nb # ((784,1),(30,1),(30,1),(10,1)) x 300
11 z_list = [[np.zeros((n, 1)) for n in layers]] * nb # ((784,1),(30,1),(30,1),(10,1)) x 300. input layerは不要だけど簡単化のためつけている。
12 a_list = [[np.zeros((n, 1)) for n in layers]] * nb # ((784,1),(30,1),(30,1),(10,1)) x 300
13
14 def SGD_train(train, ne, nb, alpha, lam = 0.0, log_interval=10, test=None):
15     """SGD for training parameters
16     If verbose is set, print progressive accuracy updates.
17     If test set is provided, routine will print accuracy on test set as learning evolves
18     """
19     for epoch in range(ne):
20         if epoch % log_interval == 0:
21             log_train_progress(train, test, epoch)
22             batch = random.sample(range(0, len(train)), nb)
23             for s, sample in enumerate(batch):
```

```

24         forward_back_prop(train[sample][0], train[sample][1], s)
25     for l in range(L-1):
26         for s in range(nb):
27             biases[l] -= (alpha/nb) * deltas[s][l+1]
28             weights[l] -= (alpha/nb) * np.dot(deltas[s][l+1], a_list[s][l].T)
29     log_train_progress(train, test, ne)
30
31 def forward_back_prop(x, y, s):
32     """Forward & back propagation for derivatives of C wrt parameters"""
33     a_list[s][0] = x
34
35     # Feedforward
36     for l in range(1, L):
37         z_list[s][l] = np.dot(weights[l-1], a_list[s][l-1]) + biases[l-1]
38         if l == L-1 and output_type == 0:
39             a_list[s][l] = softmax(z_list[s][l])
40         else:
41             a_list[s][l] = sigmoid(z_list[s][l])
42
43     # Output Error
44     if output_type == 1:
45         deltas[s][L-1] = grad_cost(a_list[s][L-1], y) * sigmoid_prime(z_list[s][L-1])
46     else:
47         deltas[s][L-1] = delta_cross_entropy(a_list[s][L-1], y)
48
49     # Back propagate
50     for l in range(L-2, -1, -1): # layerはoutputの後から遡っていく。L=4ならlayer=2,1,0
51         deltas[s][l] = np.dot(weights[l].T, deltas[s][l+1]) * sigmoid_prime(z_list[s][l])
52         # l=2の時、deltas[s][l]は(30, 1), weights[l].Tは(30, 10), deltas[s][l+1]は(10, 1), z_list[s][l]は(30, 1)
53
54 def grad_cost(a, y):
55     """gradient of cost function. Assumes C(a,y) = (a-y)^2/2"""
56     return (a - y)
57
58 def sigmoid(z, threshold=20):
59     z = np.clip(z, -threshold, threshold)
60     return 1.0 / (1.0 + np.exp(-z))
61
62 def sigmoid_prime(z):
63     return sigmoid(z) * (1.0 - sigmoid(z))
64
65 def softmax(z, overflow=False):
66     exp_z = np.exp(z)
67     if overflow:
68         exp_z = np.exp(z - np.max(z))
69     return exp_z / np.sum(exp_z)
70
71 def delta_cross_entropy(p, y):
72     return (p - y)
73
74 def forward_prop(a):
75     """forward propagation for evaluate function only"""
76     for w, b in zip(weights, biases):
77         z = np.dot(w, a) + b
78         a = sigmoid(z)
79     return a
80
81 def evaluate(data):
82     """Evaluate current model on labeled train/test data"""
83     ctr = 0
84     for x, y in data:
85         yhat = forward_prop(x)
86         ctr += yhat.argmax() == y.argmax()
87     return float(ctr) / float(len(data))
88
89 def log_train_progress(train, test, epoch):
90     """Logs training progress"""
91     acc_train = evaluate(train)
92     if test is not None:
93         acc_test = evaluate(test)
94         print("Epoch {:4d}: Train {:.10.5f}, Test {:.10.5f}".format(epoch, acc_train, acc_test))
95     else:
96         print("Epoch {:4d}: Train {:.10.5f}".format(epoch, acc_train))

```

▼ Implementation

```
1 SGD_train(train, ne, nb, alpha, lam=0.0001, log_interval=1000, test=test)
```

```
Epoch   0: Train   0.09924, Test   0.09844
Epoch 1000: Train  0.59704, Test  0.52861
```

Epoch 2000: Train	0.79672, Test	0.69508
Epoch 3000: Train	0.81192, Test	0.71068
Epoch 4000: Train	0.86715, Test	0.76110
Epoch 5000: Train	0.85314, Test	0.73870
Epoch 6000: Train	0.82513, Test	0.71789
Epoch 7000: Train	0.87155, Test	0.76991
Epoch 8000: Train	0.91236, Test	0.81192
Epoch 9000: Train	0.86515, Test	0.75350
Epoch 10000: Train	0.87315, Test	0.75870
Epoch 11000: Train	0.91477, Test	0.81032
Epoch 12000: Train	0.92837, Test	0.81353
Epoch 13000: Train	0.94278, Test	0.81913
Epoch 14000: Train	0.91957, Test	0.81232
Epoch 15000: Train	0.92677, Test	0.81353
Epoch 16000: Train	0.92237, Test	0.82313
Epoch 17000: Train	0.95598, Test	0.83713
Epoch 18000: Train	0.95078, Test	0.82593
Epoch 19000: Train	0.96319, Test	0.84674
Epoch 20000: Train	0.95238, Test	0.82753

▼ (Reference) Another Version

Update weights and biases by each image data. (not taking summation)

```
1 nl = 2
2 nh = 30
3 ne = 30
4 nb = 300
5 alpha = 0.1
6 lam = 0.0001
7 output_type = 0
```

+ コード

+ テキスト

```
1 import random
2 np.random.seed(42)
3 random.seed(42)
4
5 layers = [train[0][0].shape[0]] + [nh]*nl + [train[0][1].shape[0]]
6 L = len(layers)
7 biases = [np.random.randn(n, 1) for n in layers[1:]] # input layer以外で必要. sizeが((30,1),(30,1),(10,1))のlist of lists.
8 weights = [np.random.randn(n, m) for (m, n) in zip(layers[:-1], layers[1:])] # 長さmの縦長行列を長さnに変換する. sizeは((784,30),(30,30),(30,10)).
9 deltas = [np.zeros((n, 1)) for n in layers] # ((784,1),(30,1),(30,1),(10,1))
10 z_list = [np.zeros((n, 1)) for n in layers] # ((784,1),(30,1),(30,1),(10,1)). input layerは不要だけど簡単化のためつけている.
11 a_list = [np.zeros((n, 1)) for n in layers] # ((784,1),(30,1),(30,1),(10,1))
12
13 def SGD_train(train, ne, nb, alpha, lam = 0.0, log_interval=10, test=None):
14     """SGD for training parameters
15     If verbose is set, print progressive accuracy updates.
16     If test set is provided, routine will print accuracy on test set as learning evolves
17     """
18     for epoch in range(ne):
19         if epoch % log_interval == 0:
20             log_train_progress(train, test, epoch)
21             batch = random.sample(range(0, len(train)), nb)
22             for sample in batch:
23                 SGD_step(*train[sample], alpha, lam) # *があるのは、trainはimageとlabelの二つセットだから
24
25 def SGD_step(x, y, alpha, lam):
26     """get gradients with x, y and do SGD on weights and biases
27     Args:
28         x: single sample features. (= image data)
29         y: single sample target. (= one-hot array of label)
30     """
31     forward_back_prop(x, y)
32     for l in range(L-1):
33         biases[l] -= alpha * deltas[l+1]
34         weights[l] -= alpha * (np.dot(deltas[l+1], a_list[l].T) + lam * weights[l])
35
36 def forward_back_prop(x, y):
37     """Forward & back propagation for derivatives of C wrt parameters"""
38     a_list[0] = x
39
40     # Feedforward
41     for l in range(1, L):
42         z_list[l] = np.dot(weights[l-1], a_list[l-1]) + biases[l-1]
43         if l == L-1 and output_type == 0:
44             a_list[l] = softmax(z_list[l])
45         else:
46             a_list[l] = sigmoid(z_list[l])
47
48     # Output Error
49     if output_type == 1:
50         deltas[L-1] = grad_cost(a_list[L-1], y) * sigmoid_prime(z_list[L-1])
```

```

51     else:
52         deltas[L-1] = delta_cross_entropy(a_list[L-1], y)
53
54     # Back propagate
55     for l in range(L-2, -1, -1): # layerはoutputの後から遡っていく。L=4ならlayer=2,1,0
56         deltas[l] = np.dot(weights[l].T, deltas[l+1]) * sigmoid_prime(z_list[l])
57
58 def grad_cost(a, y):
59     """gradient of cost function. Assumes C(a,y) = (a-y)^2/2"""
60     return (a - y)
61
62 def sigmoid(z, threshold=20):
63     z = np.clip(z, -threshold, threshold)
64     return 1.0 / (1.0 + np.exp(-z))
65
66 def sigmoid_prime(z):
67     return sigmoid(z) * (1.0 - sigmoid(z))
68
69 def softmax(z, overflow=False):
70     exp_z = np.exp(z)
71     if overflow:
72         exp_z = np.exp(z - np.max(z))
73     return exp_z / np.sum(exp_z)
74
75 def delta_cross_entropy(p, y):
76     return (p - y)
77
78 def forward_prop(a):
79     """forward propagation for evaluate function only"""
80     for w, b in zip(weights, biases):
81         z = np.dot(w, a) + b
82         a = sigmoid(z)
83     return a
84
85 def evaluate(data):
86     """Evaluate current model on labeled train/test data"""
87     ctr = 0
88     for x, y in data:
89         yhat = forward_prop(x)
90         ctr += yhat.argmax() == y.argmax()
91     return float(ctr) / float(len(data))
92
93 def log_train_progress(train, test, epoch):
94     """Logs training progress"""
95     acc_train = evaluate(train)
96     if test is not None:
97         acc_test = evaluate(test)
98         print("Epoch {:4d}: Train {:.10.5f}, Test {:.10.5f}".format(epoch, acc_train, acc_test))
99     else:
100         print("Epoch {:4d}: Train {:.10.5f}".format(epoch, acc_train))

```

```

1 SGD_train(train, ne, nb, alpha, lam=0.0001, log_interval=1, test=test)

```

```

Epoch 0: Train 0.09844, Test 0.09764
Epoch 1: Train 0.47339, Test 0.41817
Epoch 2: Train 0.58824, Test 0.51220
Epoch 3: Train 0.66827, Test 0.58463
Epoch 4: Train 0.71909, Test 0.65026
Epoch 5: Train 0.74110, Test 0.64746
Epoch 6: Train 0.81232, Test 0.72829
Epoch 7: Train 0.81072, Test 0.72109
Epoch 8: Train 0.80432, Test 0.74110
Epoch 9: Train 0.83713, Test 0.74710
Epoch 10: Train 0.85714, Test 0.76591
Epoch 11: Train 0.80992, Test 0.70468
Epoch 12: Train 0.86835, Test 0.77911
Epoch 13: Train 0.85394, Test 0.77151
Epoch 14: Train 0.89116, Test 0.81593
Epoch 15: Train 0.88715, Test 0.79952
Epoch 16: Train 0.85914, Test 0.77631
Epoch 17: Train 0.87835, Test 0.77031
Epoch 18: Train 0.87875, Test 0.78151
Epoch 19: Train 0.88395, Test 0.78551
Epoch 20: Train 0.90556, Test 0.81513
Epoch 21: Train 0.90956, Test 0.81473
Epoch 22: Train 0.90396, Test 0.79832
Epoch 23: Train 0.88675, Test 0.78591
Epoch 24: Train 0.91437, Test 0.80712
Epoch 25: Train 0.89556, Test 0.81072
Epoch 26: Train 0.88635, Test 0.79032
Epoch 27: Train 0.90316, Test 0.81713
Epoch 28: Train 0.91357, Test 0.80912
Epoch 29: Train 0.92957, Test 0.82433

```